

Odysra AI

The Vision

One intelligent waiter. Every restaurant. Owned by us.

SECTION 1

The Big Idea

Odysra is building one smart “AI waiter” that can work in any restaurant. A real waiter learns the craft in one place and then moves to another carrying that experience — only the menu changes. Our AI works the same way.

The product has two parts:

The Body vs the Soul

The Body — the chatbot the customer sees (menu, ordering, QR). This can run on each restaurant’s own domain.

The Soul — the AI intelligence: how it talks, recommends, and gives insights. This stays on Odysra’s servers and is provided as a paid service.

The body is replaceable. The soul is ours. A restaurant’s chatbot only works while their Odysra account is active — so we can never be cut out.

SECTION 2

The Soul — what the AI actually is

The “soul” is three things — all written and owned by Odysra, and all reused across every restaurant:

- **Persona** — how it speaks: warm, friendly, helpful, and it knows when to suggest a drink or a dessert.
- **Recommendation logic** — how it picks dishes: by mood, taste, and time of day, mixed across categories so customers try variety.
- **Insights** — what it learns and reports back: to customers (“this pairs well with that”) and to owners (what sells, what to prep).

Because all three live in our code, every restaurant instantly gets the same experienced waiter. Only the menu is swapped in.

SECTION 3

We don’t build a brain from scratch

There are three ways to “have an AI.” We use the smart one.

- **From scratch** (like ChatGPT’s core): billions of data points, GPU clusters, millions of dollars. Not for us.
- **Use a rented brain + our data** (now): we use a ready brain (Gemini) and feed it our menus, customer history and patterns. It gets smarter as we collect data — without training anything.
- **Fine-tune our own model** (later): once we have thousands of clean conversations and orders, we teach an open model our style, so we depend less on Gemini and own more of the soul.

Path: rented brain now → our data makes it smart → fine-tune our own model at scale.

SECTION 4

The three sides

Customer side

Scan a QR or open a link → chat with the AI waiter by text or voice, in their language → “find what’s for me” gives mood-based picks → add to cart, order, pay. Faster, personal, no waiting for a waiter.

Owner side

A simple admin to manage menu and prices, plus the valuable part — insights and demand intelligence (next section). Result: fewer staff needed, bigger average orders from smart upselling, faster tables.

Odysra side

We run one system. We host the soul, hold the AI keys, improve the intelligence centrally, and meter the usage. Restaurants subscribe and get their link + QR + admin. They never touch the hard parts.

SECTION 5

Demand Intelligence — festivals & weather

The AI shouldn’t just take orders — it should think like a manager. It pulls festival / event news and weather from trusted sources (for India: IMD or a weather API for rain; a festivals calendar and local news for events) and turns them into a heads-up for the supervisor and admin.

Example heads-up

“Big temple festival near RS Puram this weekend, and a 60% chance of rain tomorrow evening — expect higher footfall. Prep extra idli batter and tiffin, and keep an extra hand on the floor.”

This makes Odysra the system owners check every morning — and it’s a strong premium feature to charge for.

SECTION 6

The data is the real asset

The moat isn’t the model — it’s the data: what people order, when, what sells together, what they avoid, which upsells work. The general patterns transfer across restaurants, while each restaurant’s private customer data stays separate and consented.

Capture it cleanly from day one — data we don’t save today, we can’t use tomorrow.

SECTION 7

The business model

- **The body is cheap or free** — it gets restaurants onboard, on their own domain.
- **The soul is metered** — this is the business.

How the money works (prepaid credits — like a phone recharge):

- Our cost per interaction = a Gemini reply + a Sarvam voice minute + a little server cost.
- We charge above that. The restaurant tops up a balance; each AI reply / voice minute deducts a little.
- Busy restaurant uses more, pays more, we earn more.
- **Premium add-on:** the insights + festival/weather demand dashboard — the highest-margin piece.

(Can also be wrapped as “Rs. X / month includes N conversations, extra billed after” for a friendlier, predictable bill.)

SECTION 8

The piece to build — the Odysra AI Gateway

Today the AI calls sit inside the app. To run this as a service, we move them to a central gateway:

- A restaurant’s chatbot calls **Odysra’s gateway** — not Gemini directly.
- The gateway checks their balance / subscription, applies our soul logic, calls Gemini / Sarvam with **our** keys, returns the answer, logs the usage, and deducts credits.
- A billing / credits system tracks balances and top-ups.

This is what lets the body live on their domain while the soul — and the money — stays with us.

SECTION 9

Roadmap

- **Now:** ship Annapoorna + Bon Bon; log every order, taste and chat cleanly; nail the persona + recommendation logic.
- **Near-term:** build the gateway + credits; add the insights + festival/weather dashboard; sell insights as premium.
- **Later:** fine-tune our own Odysra model on pooled, anonymized data; scale to many restaurants from one system.

The body gets us in the door. The soul — how it speaks, recommends, and foresees demand — is what they keep paying for. That is Odysra.