



ODYSRA

an era of good fortune

The Insight Engine

How the AI Waiter turns every conversation into insight for the owner

Briefing for Restaurant Owners

PREPARED BY

Harsith Ramprakash

Founder, Odysra

+91 74486 74496 · odysra.official@gmail.com

June 2026

Confidential — Odysra capability briefing.

1. Why this matters

Every restaurant already knows **what** sold. Almost none know **why** — why one dish flies and another stalls, what a guest actually wanted, or how they felt about the plate once it arrived. Feedback forms catch only a handful of guests, and even then only about “service” or “taste” in general. The Odysra AI Waiter changes this: because guests **talk to it** before they order, every conversation can be turned into structured insight — and the owner gets a read on **every dish**, not just the few that ever get rated.

2. How we capture data from every conversation

The waiter quietly labels each interaction *as it talks* — no extra step for the guest, no separate system. For every AI-assisted order it records, in plain structured form:

- **What the guest asked for** — the intent in their words: “something light”, “spicy”, “filling”, “for kids”, “under Rs.150”, “what’s famous”.
- **What the AI recommended** — which dishes it suggested, and in what order.
- **What happened next** — viewed → added → ordered, or dropped; and *how* they chose (took the suggestion, asked by name, or browsed).
- **The context** — language, time of day, dine-in vs takeaway, party size, branch.
- **How they felt** — the words they used (“too spicy”, “loved it”), plus the explicit topic rating at the end.

One conversation, captured: A guest asks for “*something filling, not too spicy.*” The AI offers Biryani, Meals and Pulav. The guest opens Biryani and adds it. Stored: intent = *filling / mild*; recommended = [biryani, meals, pulav]; chosen = **biryani** via *accepted recommendation*; lunch, Telugu, dine-in. Multiply this by thousands of conversations and clear patterns appear.

3. From conversations to trends

Those labelled conversations roll up — by the database, exactly and cheaply — into summary tables the owner can act on. A few examples:

- **Intent → dish** — for “filling” requests, Biryani is offered 70% of the time and accepted 55%.
- **Recommendation acceptance** — which dishes guests take when suggested (Biryani 55%, Dish Y 8%).
- **View → order funnel** — Dish Y: viewed 300, added 40, ordered 12 — the drop is at ‘add to order’, a price signal.
- **Why-picked mix** — Biryani: 60% accepted the suggestion, 30% asked by name, 10% browsed.
- **Per-dish sentiment, pairings, and audience** — who orders it, when, with what, and what they say about it.

4. Feedback on every dish — without a form

This is the part owners have never had. Guests rarely rate each dish — but their **behaviour and words** reveal how a dish is really doing. From the conversation and order data, the engine reads a satisfaction signal for **every item on the menu**, not just the few that get a star rating:

- **Re-order & repeat rate** — do the same guests come back to it? The strongest vote of all.
- **Accepted-and-kept vs viewed-and-skipped** — a dish people look at but never add is quietly failing.
- **Where it drops** — dropping after seeing the price vs after seeing the description points to different fixes.
- **In-chat words** — “too oily”, “less quantity”, “best dosa here” — captured per dish, in any language.
- **Recommendation success** — a dish the AI keeps suggesting that nobody accepts is telling you something.

So instead of a vague “food: 4 stars,” the owner sees, dish by dish, whether guests love it, ignore it, or quietly reject it — and a likely reason.

5. How the owner sees it — reasons, then action

The AI reads only these small summary tables and writes the finding in plain English — the reason, then what to do:

Biryani sells because it is the default answer for ‘filling / mild’ requests and 55% accept it; it peaks 12–2pm with Telugu and Hindi guests, usually with a cold drink.

Dish Y lags because it is almost never recommended and 70% who view it don’t order — most drop at the price step. **Do this:** offer Y when guests ask for ‘light’, trial it in a combo, and feature it at lunch — then watch next week’s numbers.

Those actions feed back into the waiter’s recommendations, the change is measured, and the engine refines itself — a loop that keeps sharpening. Owners read it per branch in the supervisor view; the brand sees every branch together in the admin dashboard.

6. Trustworthy, and inexpensive

- **The database does the counting** — exact, instant, and effectively free; the AI only *interprets* small summaries, so running cost stays a few thousand rupees a month even at depth.
- **Reasons are proposals, not verdicts** — the engine surfaces the *likely* cause and the owner confirms it; it never hides the difference between a pattern and a guess.
- **Privacy-light** — insight comes from orders and conversations, not personal data; only a phone number (for the bill) is ever asked of the guest.

In short: the AI Waiter is not just taking orders — it is listening, at every table, in every language, and turning what it hears into the clearest picture a restaurant has ever had of why its menu performs the way it does.